Analysis of SARSA, Q-Learning and Monte-Carlo Techniques on Taxi Environment

Course Project - Introduction to Reinforcement Learning

Laiba Jamil
Muhammad Zain Yousuf
Hamad Abdul Razzaq

Problem
Taxi Environment

- 2D Grid $\rightarrow 5 \times 5$
- Driver needs to pick and drop the passenger
- Passenger can be in 4 locations
- Destination can also be 4 locations
- We will model this problem as MDP

Modeling the Problem as MDP

- **State Space**
  - 4-tuple

$$s = (d_x, d_y, p_p, p_d)$$

  - d_x: x-position of the driver
  - d_y: y-position of the driver
  - p_p: Current position of Passenger
  - p_d: Drop Off Location

Modeling the Problem as MDP

- **Action Space**
  - Driver can move in 4 directions
  - Driver can pick/drop the passenger

$$A = \{ \text{Left, Right, Up, Down, Pick, Drop} \}$$

Modeling the Problem as MDP

- **Reward Shaping**

  - Driver moves in either direction $\rightarrow r = -1$
  - Driver drops passenger in wrong location $\rightarrow r = -10$
  - Driver drops passenger in correct location $\rightarrow r = 20$

Hyper-Parameters

- $N = 1000$ Episodes
- Max Episode Length = 50000
- $\gamma = 0.7$ (Discount Factor)
- $\alpha = 0.1$ (Learning Rate)
- $\epsilon = 0.1$ (Exploration)

Applying Q-Learning Technique

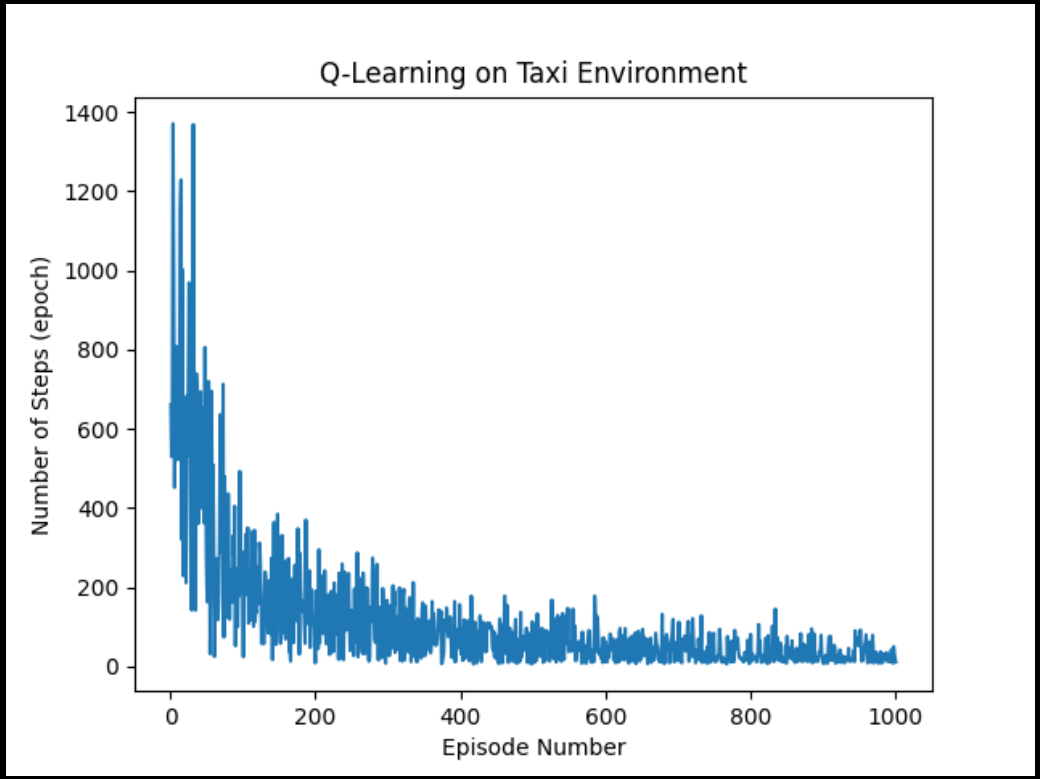

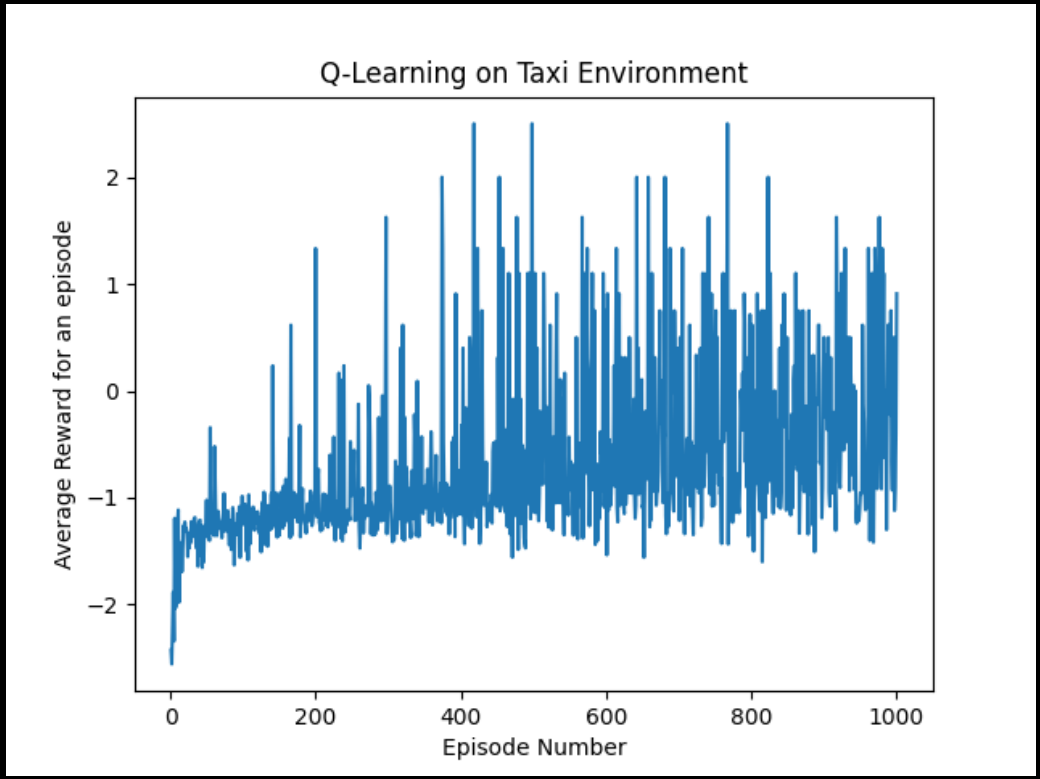

Average Reward = 0.13

Applying SARSA Technique

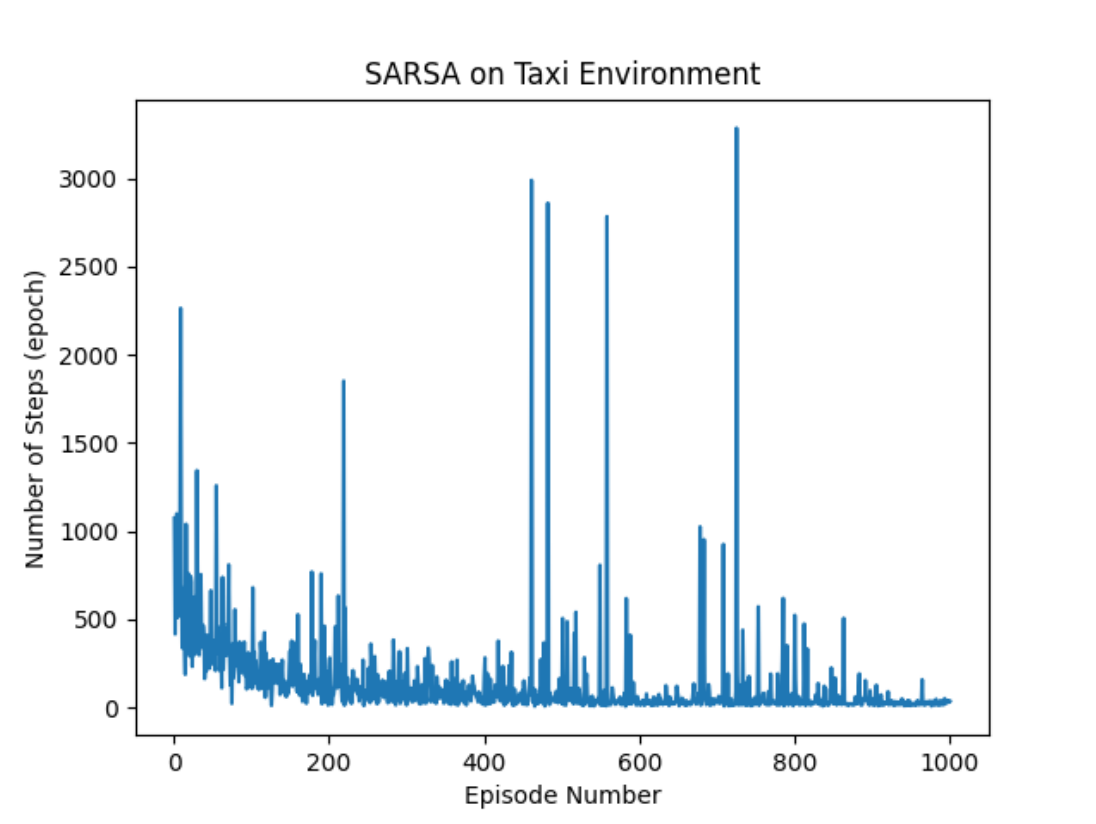

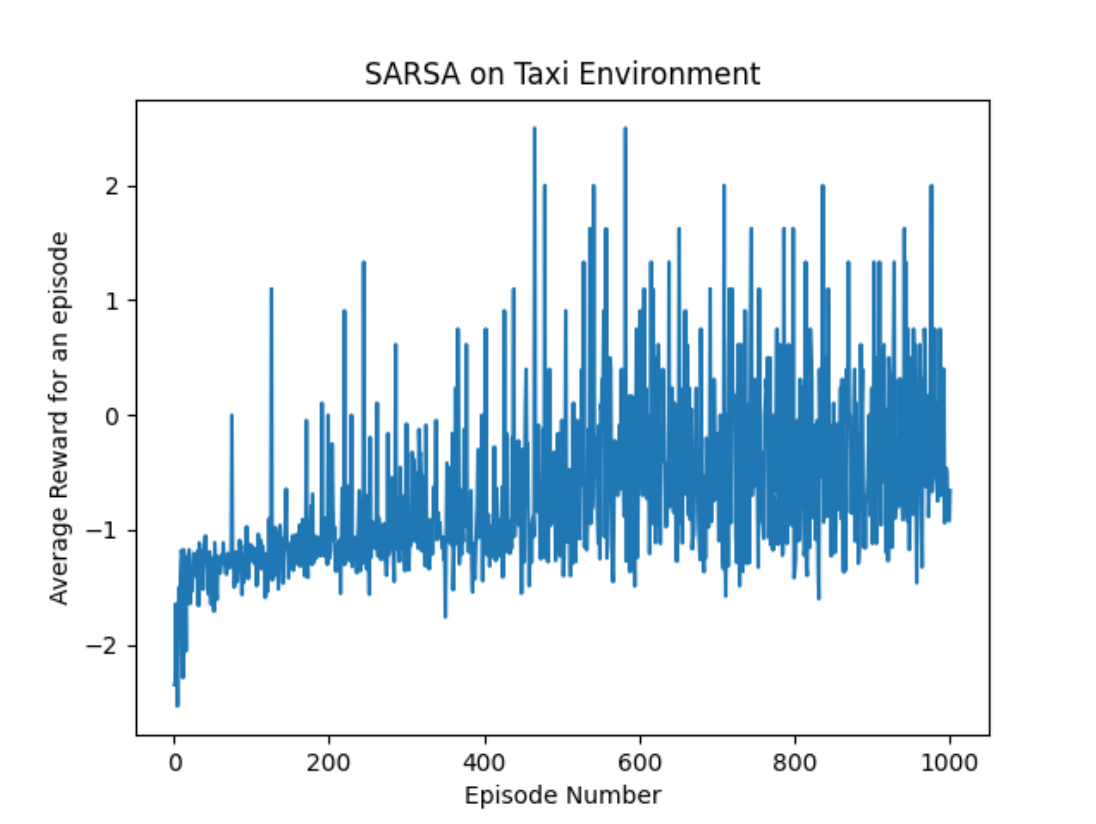

Average Reward = −0.09

Applying Monte-Carlo Technique

- Every Visit

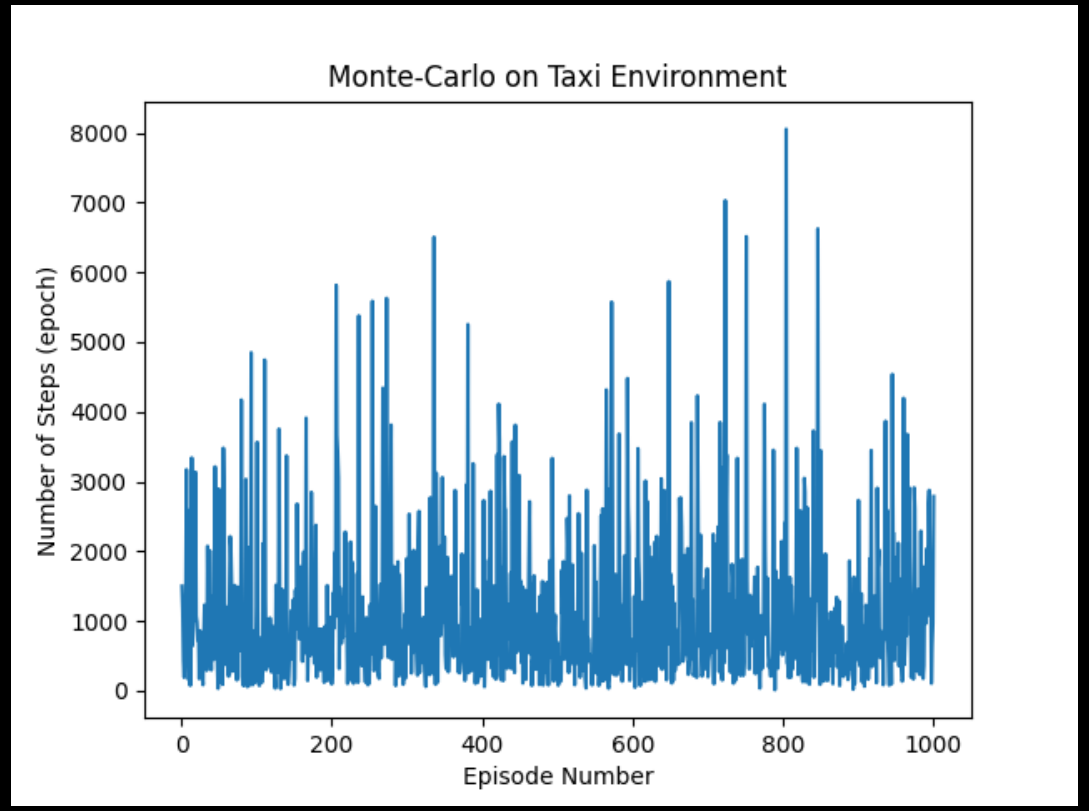

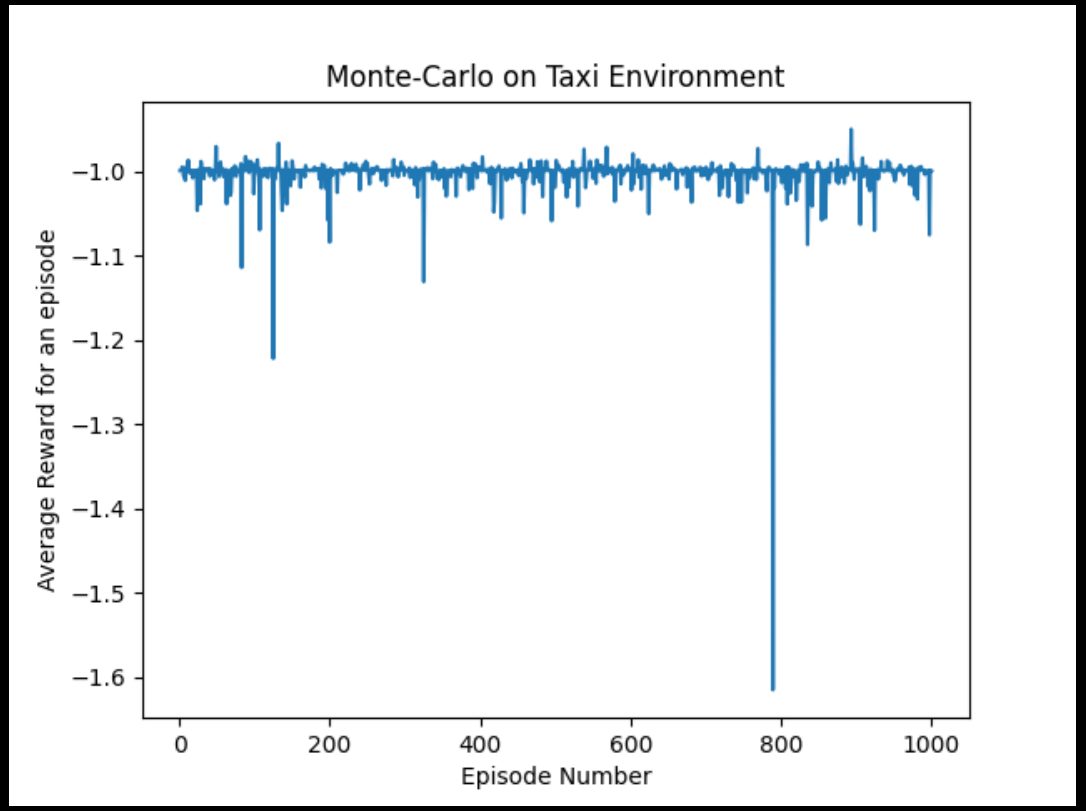

Average Reward = −1.002

Analysis / Comparison

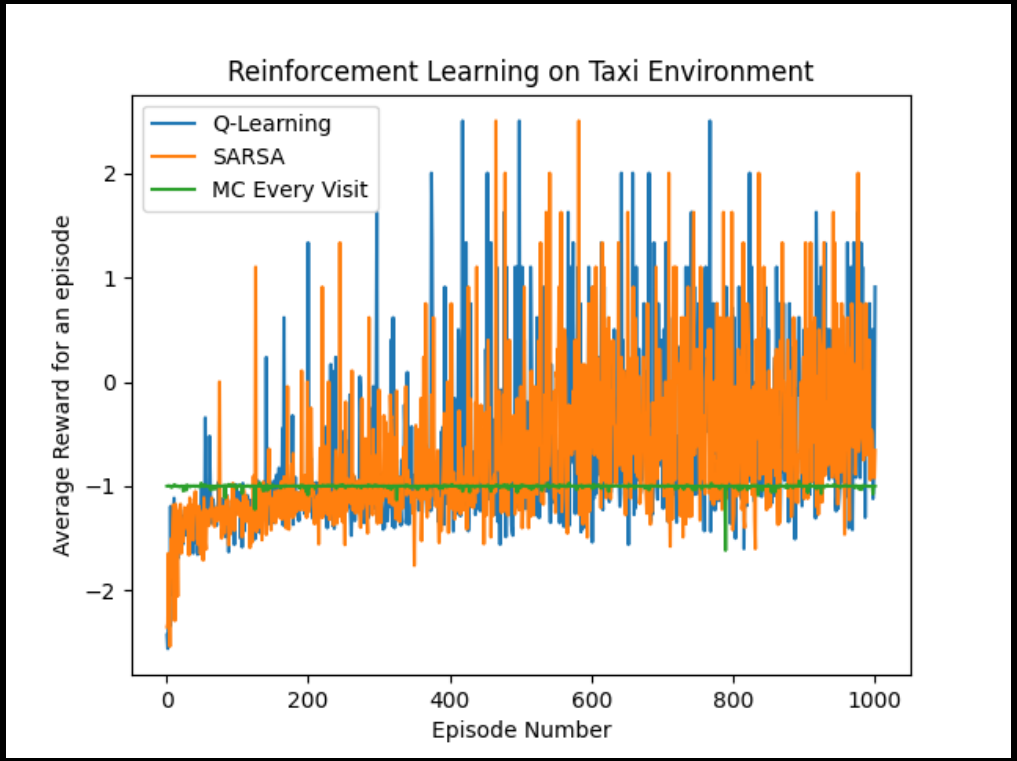

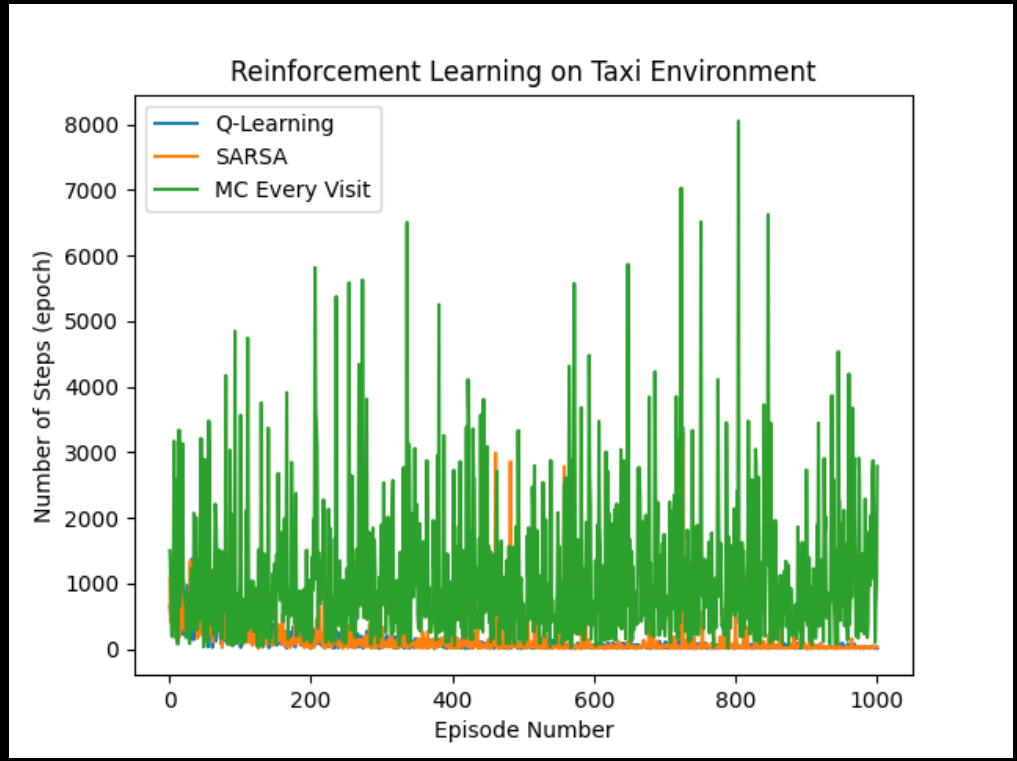

Q-Learning Outperforms!

Thank you!

References:

- [1] https://towardsdatascience.com/solving-the-taxi-environment-with-q-learning-a-tutorial-c76c22fc5d8f
- [2] R. S. Sutton and A. G. Barto, Reinforcement Learning: An Introduction. Cambridge, MA: The MIT Press, 2020.